

CDDC 2019 Qualifiers

Challenge Writeup

Team: S4F3 S0L1D5

Category: Uni/Poly

June 2019

OSINT_RED

1 [R-0] Everyone <3 Fan Mail

This series of challenges were made slightly more tedious than necessary, because unfortunately there was a real company called 'LightSpeed' that diluted the search results.

The main objective of this challenge was to find the email of the site administrator for www.lightspeedcorp.global, and send him an email.

1.1 Solution

In theory, all we would have needed to do was perform a WHOIS lookup on the domain to find the email address. Simple, right?

Nope.

Depending on which WHOIS search engine is used, certain bits of information about the registrant can be missing if the domain was registered with WHOISGuard, which hides certain identifiable information from WHOIS lookups.

```
Registrant Name: Luther Torvalds
Registrant Organization: LightSpeedCorp
Registrant Street: Add 1 Add 2
Registrant City: Singapore
Registrant State/Province: NA
Registrant Postal Code: 123456
Registrant Country: SG
Registrant Phone: +65.62353535
Registrant Phone Ext: 213
Registrant Fax:
Registrant Fax Ext:
Registrant Email: Luther.Torvalds@outlook.com
```

Figure 1.1: All the information is visible if WHOISGuard is ignored

However, it turns out that the guard isn't perfect, and all we had to do was find a WHOIS search engine that ignores the guard. After some digging, we found one, exposing all the juicy details of the registrant.

After getting the email address, we just had to send an email, and we were replied with the flag.

1.2 Flag

The flag for this challenge was `$CDDC19${IS_IT_I_AM_FAMOUS_NAO}`.

2 [R-1] Travel to the Past

The challenge heavily hinted that we would need a way to look at an older version of the website to find the flag, and that's exactly what we did.

2.1 Solution

There was a single snapshot of the website on the Wayback Machine¹ on the 20th of May, and browsing that snapshot revealed the flag on the homepage of a blog:

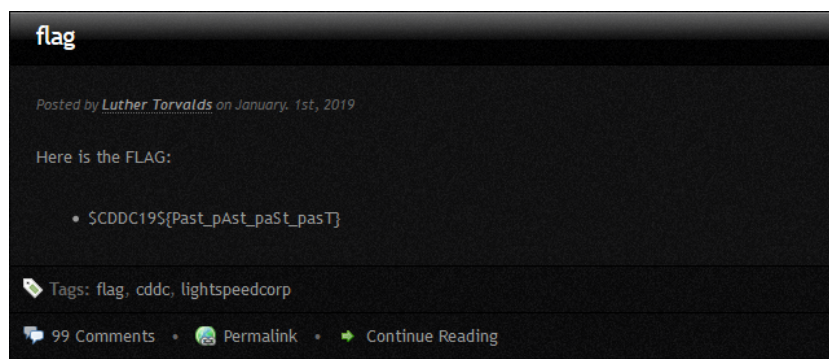


Figure 2.1: The flag is clearly visible

2.2 Flag

The flag for this challenge was `$CDDC19${Past_pAst_paSt_pasT}`.

[1] <https://archive.org/web/>

3 [R-2] I'm Sho Done With This

3.1 Solution

The hint was in the question title itself — turns out that there is a search engine called ‘ShoDan’², which claims to be the world’s first search engine for internet connected devices.

Sure enough, searching for our favourite company ‘lightspeedcorp’ yielded the following results:

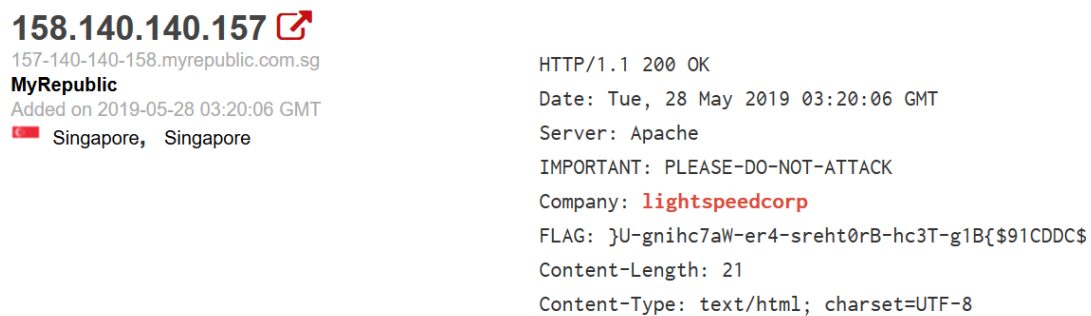


Figure 3.1: The flag is reversed, for some reason

A trivial string-reverse later, the flag is obtained.

3.2 Flag

The flag for this challenge was \$CDDC19\${B1g-T3ch-Br0thers-4re-Wa7ching-U}.

[2] <https://www.shodan.io>

Programming

4 Count 1: Baby

This was a fairly simple code golfing challenge; inspection of the provided `count1-baby.py` file gave the character limit as 53 characters. The starting code is given as this:

```
1  #include <stdio.h>
2
3  int main(int argc, char *argv[])
4  {
5      int i;
6
7      for( i = 1 ; i < 10000 ; i++ )
8      {
9          printf("%d,", i);
10     }
11
12     return 0;
13 }
```

4.1 Solution

Compiling and running the program shows that the expected output of the minified program are the numbers 1 to 9999, separated by commas (with a trailing comma after 9999).

Besides the obvious steps like removing whitespace and newlines, knowledge of the C standard and ignorance of compiler warnings can yield the following transformations:

- `printf` is an implicitly defined standard library function; `stdio.h` can be removed
- Functions do not require a type specifier; `int main` can be shortened to `main`
- `main` does not need to take arguments, leaving `main()`
- The verifier rejects spaces, so `int i` has to be changed to `int(i)` (which is valid)
- `main` does not need to explicitly return 0

This yields the final program, at 53 characters long:

```
1 | main(){int(i);for(i=1;i<10000;i++){printf("%d",i);}}
```

Listing 4.1: The solution for Count 1: Baby

4.2 Flag

The flag for this challenge was `$CDDC19${Count2_is_waiting_Please_enjoy!}`.

5 Count 2: Wildness

The required output for this challenge is the same as the previous challenge (Count 1: Baby), but the restrictions have been tightened. This time, OP's friend 'doesn't like to go wild', or, in other words, the following characters cannot be used in the source code: `<== wlld ==>`. Additionally, the character limit has been reduced to 41.

5.1 Solution

Starting from the code for Count 1 (Listing 4.1), we can apply some more non-obvious tricks:

- Variables don't need a type specifier, and will default to `int`
- Variables with static storage duration will be initialised to 0 (obviating the need for `=`)
- `!=` is equivalent to a bitwise XOR (^) for integers (removing `<`)
- Taking advantage of `i++` semantics can replace `1e4` with `9999` (since `1` can't be used)

This yields the final program, again at exactly 41 characters:

```
1 | i;main(){for(;i++^9999;printf("%d",i));}
```

Listing 5.1: The solution for Count 2: Wildness

5.2 Flag

The flag for this challenge was `$CDDC19${This_really_helps_m3_a_lot}`

6 Count 4: Madness - Filter

For this challenge, the character limit is slightly relaxed to 44 characters, but OP's friend has a faulty keyboard now, and all lowercase characters except those in 'mad printf' ('a', 'd', 'f', 'i', 'm', 'n', 'p', 'r', and 't') cannot be used.

6.1 Solution

Given that all the looping constructs (for, while, and even goto) cannot be used due to having illegal characters, the only logical solution is recursion.

- It is not illegal to call main recursively
- When called with no arguments, argc is 1; so main(i) will initialise i to 1
- 1E4 is equally as valid as 1e4, which is identical to 10000
- The first arm of the ternary operator can be omitted: x?:y is valid.

We came up with two independent solutions to this problem:

```
1 | main(i){printf("%d,",i++);if(i<1E4)main(i);}
```

Listing 6.1: The solution for Count 2: Wildness

```
1 | i;main(){printf("%d,",++i);i>=9999?:main();}
```

Listing 6.2: An alternative solution for Count 2: Wildness

In both cases, the code was exactly 44 characters long.

6.2 Flag

The flag was `$CDDC19${Main_might_be_just_a_function_but_it_is_really_special!}`.

Reverse Engineering

7 LSCVM: Immaculate Invasion

Presented with a login prompt, the logical conclusion was that the flag would be accessible upon successfully logging into the server.

7.1 Solution

The tool of choice for this challenge was *Cutter*³, an open-source reverse-engineering framework which performs disassembly, function analysis, and function/value renaming, among other things.

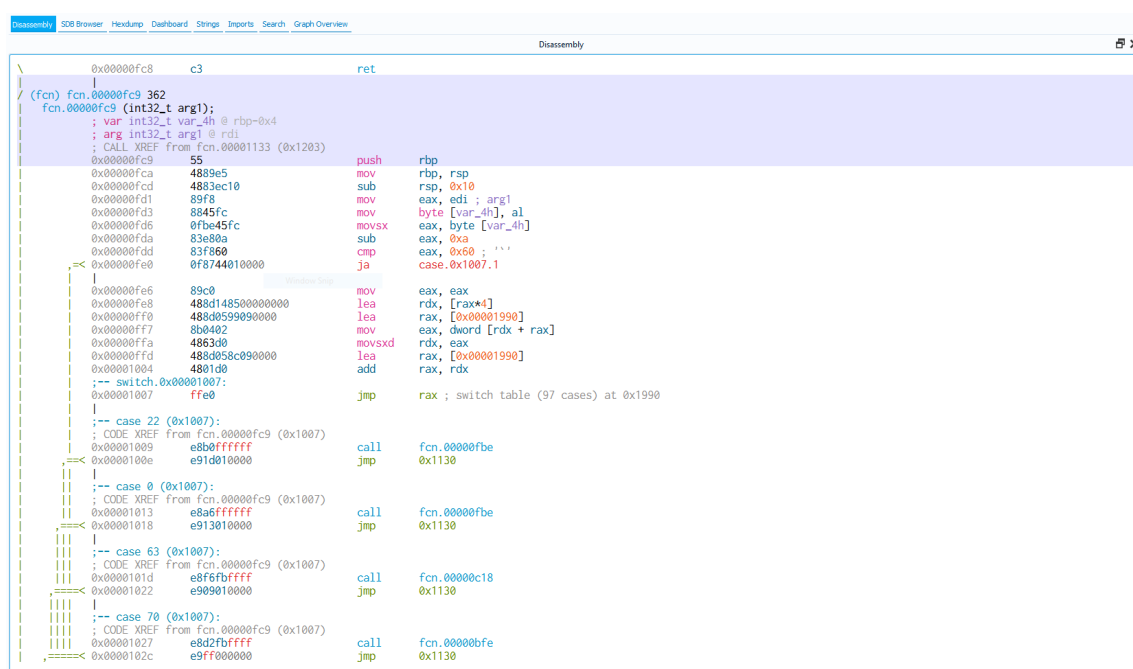


Figure 7.1: A screenshot of Cutter

The first hurdle was getting the program to run beyond its error message of [-] Flag file open error; upon inspection of the disassembly, the program tries to open a file named ‘flag’, which presumably contains the flag on the server-side — \$ touch flag convinced the program to start.

[3] <https://github.com/radareorg/cutter>

Given that the program asks for a login ID, the next step was to check for strings and `strcmps` in the code, and one was indeed found:

```
call    fcn.00001133
lea     rsi, str.lsc_user ; 0x2065 ; "lsc_user" ; const char *s2
lea     rdi, [0x00203140] ; 0x203140 ; const char *s1
call    sym.imp.strcmp ; int strcmp(const char *s1, const char *s2)
test    eax, eax
je      0x14b3
```

Figure 7.2: A suspicious `strcmp`

Of course, no challenge is so easy, especially one that sat at above 950 points more than 24 hours after the qualifiers started. On further inspection, the operand of `strcmp` appeared to be the output of a call to `fcn.00001133`, which itself took some very long and suspicious strings as input: `eiMaAghMcAgjcMMdAgjcMMaAgjcMMAhhcMMdAijMaAcfMPPPPPPPP`.

Another discovery was this instruction: `cmp dword [var_1174h], 2`, where `var_1174h` contained the value of `argc`. Passing a second argument in the invocation (eg. `$./lscvm-ii x`) revealed a very useful debugging mode that printed what appeared to be a stack.

Following the call chain eventually led to the jump table for opcode dispatch, pictured in Figure 7.1. Some quick analysis revealed that this was a indeed stack-based virtual machine (*thankfully*). Combined with the debugging output, and the fact that the text output (the banner, login prompt, etc.) was printed using the VM instead of say `printf`, some basic opcodes were deciphered:

| Opcode | Pops | Pushes | Description |
|---------------|------|--------|------------------------------|
| u to u | – | 0 to 9 | constant |
| A | a, b | a + b | add |
| M | a, b | a * b | multiply |
| P | a | – | <code>printf("%c", a)</code> |

Piecing things together, it was inferred that we were most likely supposed to input a *program* when asked for the user ID, which would then print `lsc_user` as required. Given that ASCII values for the lowercase alphabets start at 97, a small program was quickly written that would take in a string and output opcodes that would print it.

After trying it out, however, we were quickly met with disappointment, as we were faced with our old friend `[-] Wrong id.`

After more digging, the user id was expected to be at the memory address 0x203140; using the renaming feature of Cutter, we could make it easier to spot:

```
call    fcn.00001133
lea     rsi, str.lsc_user ; 0x2065 ; "lsc_user" ; const char *s2
lea     rdi, var.mem_base.203140 ; 0x203140 ; const char *s1
call    sym.imp.strcmp ; int strcmp(const char *s1, const char *s2)
test    eax, eax
je      0x14b3
```

Figure 7.3: [0x203140] after being renamed into var.mem_base.203140

Scrolling through the opcode list while looking for the address (now renamed) yielded this very interesting function, mapped to opcode K:

```
(fcn) fcn.00000e22 86
fcn.00000e22 ();
; var signed int var_4h @ rbp-0x4
; CALL XREF from fcn.00000fc9 (0x1107)
0x00000e22      55                push    rbp
0x00000e23      4889e5            mov     rbp, rsp
0x00000e26      4883ec10          sub     rsp, 0x10
0x00000e2a      e89cfdffffff      call    fcn.00000bcb
0x00000e2f      8945fc            mov     dword [var_4h], eax
0x00000e32      817dfc00ffffff     cmp     dword [var_4h], 0xffffffff
0x00000e39      7c09              jl      0xe44
; <=
|
| 0x00000e3b      817dfc003f0000     cmp     dword [var_4h], 0x3fff
| 0x00000e42      7e16              jle     0xe5a
|
| ; CODE XREF from fcn.00000e22 (0xe39)
| -> 0x00000e44      8b45fc            mov     eax, dword [var_4h]
| 0x00000e47      89c6              mov     esi, eax
| 0x00000e49      488d3de80a0000     lea     rdi, str.memory_write_access_violation_d ; 0x1938
| 0x00000e50      b800000000        mov     eax, 0
| 0x00000e55      e876fbffffff      call    sym.imp.printf ; int printf(const char *format)
|
| ; CODE XREF from fcn.00000e22 (0xe42)
| -> 0x00000e5a      e86cfdffffff      call    fcn.00000bcb
| 0x00000e5f      89c1              mov     ecx, eax
| 0x00000e61      8b45fc            mov     eax, dword [var_4h]
| 0x00000e64      4863d0            movsxd  rdx, eax
| 0x00000e67      488d05d2222000     lea     rax, var.mem_base.203140 ; 0x203140
| 0x00000e6e      880c02            mov     byte [rdx + rax], cl
| 0x00000e71      b800000000        mov     eax, 0
| 0x00000e76      c9                leave   eax
| 0x00000e77      c3                ret
```

Figure 7.4: A function to ~~surpass-metal-gear~~ write to memory?

There was also an accompanying opcode E, which read a value from memory. The memory buffer appeared to be an array of 32-bit values, and opcodes K and E took an index into this array.

Slight modifications were made to lscvm-deasciinator — keeping track of an offset and replacing ‘P’ with ‘K’ — which would generate opcodes to write a string into a given offset in memory, yielding lscvm-memoryinator (Listing 1.1).

The password was also in plaintext alongside the username: hi_darkspeed-corp!. Finally, then, the challenge could be solved:

```
$ ./lscvm-memoryinator
string: lsc_user
address: 0

ggdMMaKfgAfMcMfAbKjjMjAjAcKgfdMMfAdKfgAfMcMhAeKfgAfMcMfAfKcf
McMfMMbAgKfgAfMcMeAhK

string: hi_darkspeed-corp!
address: 0

jeAiMaKhdfMMbKgfdMMfAcKcfMcMfMMdKjjMjAhAeKfgAfMcMeAfKhdfMMcAg
KfgAfMcMfAhKfgAfMcMcAiKcfMcMfMMbAjKcfMcMfMMbAcfMKcfMcMfMMbcfMMb
AKfddMMbcfMMcAKjjMjAjAbcfMMdAKfgAfMcMbAbcfMMeAKfgAfMcMeAbcfM
MfAKfgAfMcMcAbcfMMgAKfgAdMbcfMMhAK
^C

$ nc lscvm-ii.cddc19q.ctf.sg 9001

=== Welcome to LSCVM(LightSpeed Corp Virtual Machine) ===
ID : ggd...AhK
Password : jeA...hAK

Login Successful! $CDDC19${IcY_GrE37ings_Fr0M_LigHT5pEeDC0Rp}

lsc_user, Good Bye!
```

7.2 Flag

The flag for this challenge was `$CDDC19${IcY_GrE37ings_Fr0M_LigHT5pEeDC0Rp}`.

7.3 Further Analysis

While we were unable to actually come up with a quine to submit for the next LSCVM challenge (Quintessential Harlequin), we continued to take apart the VM (on the advice that it would be used again in the Finals!), since solving lscvm-ii did not require all of the opcodes (far from it).

Again, the renaming feature of Cutter was extremely helpful, and we were able to discover the address of the program counter, the base address of the stack, and (in the case of lscvm-qh), the mirror of stdout. Each buffer appears to be identically constructed with a fixed size of 0x4e200 bytes (0x13880 32-bit words), and the number of items stored after the last element (ie. at offset 0x4e200 from the base address).

We (eventually) managed to decipher all of the opcodes and their purpose:

| Opcode | Pops | Pushes | Description |
|---------------|-----------|--------|---|
| a to j | – | 0 to 9 | constant |
| A | a, b | a + b | add |
| B | – | | stop execution immediately |
| C | x | – | call (jump to absolute instruction x) |
| D | x | – | pop (drop) |
| E | addr | value | read memory from addr |
| F | ofs | value | fetch from stack (ofs elms below top) |
| G | ofs | – | relative jump forward |
| H | ofs | value | same as F , but removes the element |
| I | x | – | printf("%d", x) |
| J | a, b | cmp | –1 if $a < b$, 0 if $a = b$, 1 if $a > b$ |
| K | val, addr | – | writes val to memory at addr |
| M | a, b | a * b | multiply |
| P | x | – | printf("%c", x) |
| R | – | – | return |
| S | a, b | a - b | subtract |
| V | a, b | a / b | integer divide |
| Z | cond, ofs | – | jump (relative) if cond is 0 |

Of note are the **C** and **R** opcodes, which *call* and *return* respectively. There is another array which is only accessed by these instructions that functions as a *callstack*. **C** pushes the current program counter (*PC*) to this callstack, and **R** pops the return address from the callstack, and sets *PC* to it, moving execution back to the callsite.

Miscellaneous

8 Polyglot

What to do, when presented with languages that we can't understand?

8.1 Solution

The obvious solution was to put each sentence into Google Translate⁴; all 10 sentences were some variation of 'the first letter of this language is the flag'. Coupled with the input format, given as [01][02][03][04][05][06]&[07][08][09][10]!, we were able to decode the flag.

| Number | Language | Sentence |
|--------|------------|---|
| 01 | Hindi | इस भाषा का पहला चरित्र झंडा बनाता है। |
| 02 | Indonesian | Karakter pertama bahasa ini yang mengibarkan bendera. |
| 03 | Chinese | 这种语言的第一个字符构成了旗帜。 |
| 04 | Dutch | Het eerste teken van deze taal vormt de vlag. |
| 05 | Danish | Det første tegn på dette sprog udgør flag. |
| 06 | Catalan | El primer caràcter d'aquest idioma constitueix la bandera. |
| 07 | Norwegian | Det første tegnet av dette språket utgjør flagget. |
| 08 | Spanish | El primer carácter de este lenguaje lo constituye la bandera. |
| 09 | Hmong | Thawj qhov cim ntawm hom lus no ua rau tus chij. |
| 10 | Croatian | Prvi znak ovog jezika čini zastavu. |

Assembled, the message was HI~CDDC&NSHC; it helped that there was a coherent message to verify that Google Translate didn't misdetect any of the languages.

8.2 Flag

The flag for this challenge was `$CDDC19${HI~CDDC&NSHC}`.

[4] <https://translate.google.com>

9 Do You Fancy Numbers?

We are presented with a picture of what are apparently numbers. On first guess they appear to be Chinese in nature.

9.1 Solution

Indeed, a quick trawl of Wikipedia led us to Suzhou numerals⁵, which contained everything we needed to decode the message.

After substituting them for arabic numerals, we get the following string of numbers: 36 67 68 68 67 49 57 36 123 53 48 95 121 48 117 95 102 52 78 99 89 95 102 108 48 87 51 114 95 78 117 77 98 51 82 53 125.

Putting them through an ASCII decoder yields the following Base64 encoded string: JCBdIEQgRCB DIDEg0SAkIHsgNSAwIF8geSAwIHUgXyBmIDQgTiBjIFkgXyBmIGwgMCBXIDMgcjBfIE4gdSBNIGIgMy BSIDUgfQ==

Finally, putting that through a Base64 decoder yields the flag (with some spaces).

9.2 Flag

The flag for this challenge was `$CDDC19${50_y0u_f4NcY_fl0W3r_NuMb3R5}`.

[5] https://en.wikipedia.org/wiki/Suzhou_numerals

10 Super Strong TeleVision

With the obvious hints in the challenge statement, we got to work decoding. SSTV, or Slow Scan Television, is a picture transmission method used mainly by amateur radio operators to transmit and receive static pictures via radio in monochrome or colour^[6]. It can also be used to hide easter^[7] eggs^[8]...

10.1 Solution

The solution was rather straightforward. Using *PulseAudio* as a link between *VLC* and *QSSTV*^[9] (an open-source Linux SSTV application), we were able to decode the SSTV image rather quickly:



Figure 10.1: QSSTV decoding the image; its frequency spectrum can be seen on the right

In this case, PulseAudio acted as the pipe that redirects the .wav file as an input into the QSSTV decoder:




Figure 10.2: Data Pipeline Setup

[6] en.wikipedia.org/wiki/Slow-scan_television

[7] https://wiki.kerbalspaceprogram.com/wiki/List_of_easter_eggs

[8] https://half-life.fandom.com/wiki/Portal_ARG

[9] <http://users.telenet.be/on4qz/qsstv/index.html>

This works by enabling and making use of the default null-sink provided with every PulseAudio installation:

```
$ pactl load-module module-null-sink sink_name=virtual-cable  
$ pavucontrol
```

and making QSSTV record from the null sink:

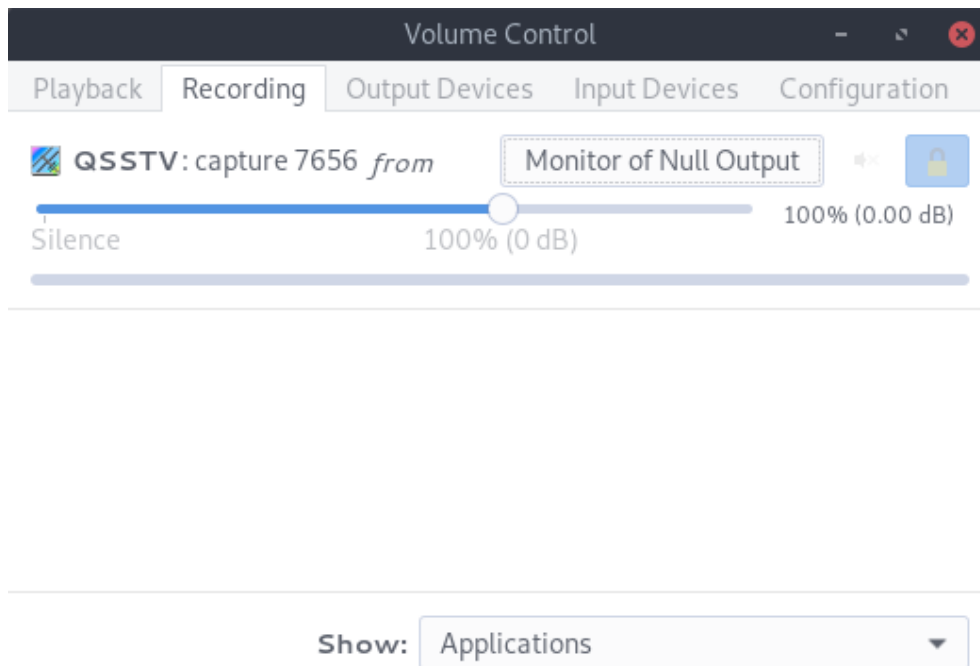


Figure 10.3: QSSTV – Record from the null-sink

Once everything is set up, we simply start recording on QSSTV, use VLC to playback the .wav file to the null-sink, and watch the magic unfold.

10.2 Flag

After some suspense as the SSTV audio played back in real time, the image was fully decoded, revealing the flag, which was `$CDDC19${Light$peedCorp-$$TV}`.



Figure 10.4: Decoded SSTV Image

Appendix

1 Code Listings

1.1 lscvm-memoryinator

```
1 // lscvm-memoryinator.cpp
2
3 #include <stdio.h>
4
5 #include <map>
6 #include <vector>
7 #include <string>
8 #include <iostream>
9
10 std::map<char, std::string> lookup;
11
12 std::string createNumber(int num)
13 {
14     if(num < 10)
15     {
16         return std::string(1, (char) (num + 'a'));
17     }
18     else if(num == 10)
19     {
20         return "cfM";
21     }
22     else if((num >= 'a' && num <= 'z') || num == '-'
23             || num == '_' || num == '!')
24     {
25         return lookup[num];
26     }
27     else
28     {
29         std::string ret;
30
31         int x = num / 10;
32         ret = createNumber(x) + createNumber(10) + "M";
```

```

33
34         int y = num % 10;
35         ret += createNumber(y) + "A";
36
37         return ret;
38     }
39 }
40
41
42 int main()
43 {
44     lookup['a'] = "jjMjAhA";
45     lookup['b'] = "jjMjAiA";
46     lookup['c'] = "jjMjAjA";
47     lookup['d'] = "cfMc fMM";
48     lookup['e'] = "cfMc fMMbA";
49     lookup['f'] = "jiAc dMM";
50     lookup['g'] = "jiAc dMMbA";
51     lookup['h'] = "jeAiM";
52     lookup['i'] = "hd fMM";
53     lookup['j'] = "hd fMMbA";
54     lookup['k'] = "hd fMMcA";
55     lookup['l'] = "gg dMM";
56     lookup['m'] = "gg dMMbA";
57     lookup['n'] = "fgAfMcM";
58     lookup['o'] = "fgAfMcMbA";
59     lookup['p'] = "fgAfMcMcA";
60     lookup['q'] = "fgAfMcMdA";
61     lookup['r'] = "fgAfMcMeA";
62     lookup['s'] = "fgAfMcMfA";
63     lookup['t'] = "fgAfMcMgA";
64     lookup['u'] = "fgAfMcMhA";
65     lookup['v'] = "fgAfMcMiA";
66     lookup['w'] = "fgAfMcMjA";
67     lookup['x'] = "gcfcMMM";
68     lookup['y'] = "fgAfgAM";
69     lookup['z'] = "fgAfgAMbA";
70     lookup['_'] = "gfdMMfA";
71     lookup['!'] = "fgAdM";
72     lookup['-'] = "fddMM";
73
74
75     std::string input;
76
77     while(true)
78     {

```

```

79     printf("string: ");
80     std::getline(std::cin, input);
81
82     printf("address: ");
83     std::string ofs;
84     std::getline(std::cin, ofs);
85
86     int offset = 0;
87     if(!ofs.empty()) offset = std::stol(ofs);
88
89     std::string output;
90     for(size_t i = 0; i < input.size(); i++)
91     {
92         // K writes to memory. format: K [value] [address]
93         output += createNumber(input[i]); // this is the value
94         output += createNumber(offset);   // this is the offset
95         output += "K";
96
97         offset++;
98     }
99
100     printf("\n%s\n", output.c_str());
101 }
102 }
```

Listing 1.1: lscvm-memoryinator.cpp