

CDDC 2019 Qualifiers

Challenge Writeup

Team: S4F3 S0L1D5

Category: Uni/Poly

June 2019

OSINT_RED

1 [R-0] Everyone <3 Fan Mail

This series of challenges were made slightly more tedious than necessary, because unfortunately there was a real company called 'LightSpeed' that diluted the search results.

The main objective of this challenge was to find the email of the site administrator for www.lightspeedcorp.global, and send him an email.

1.1 Solution

In theory, all we would have needed to do was perform a WHOIS lookup on the domain to find the email address. Simple, right?

Nope.

Depending on which WHOIS search engine is used, certain bits of information about the registrant can be missing if the domain was registered with WHOISGuard, which hides certain identifiable information from WHOIS lookups.

```
Registrant Name: Luther Torvalds
Registrant Organization: LightSpeedCorp
Registrant Street: Add 1 Add 2
Registrant City: Singapore
Registrant State/Province: NA
Registrant Postal Code: 123456
Registrant Country: SG
Registrant Phone: +65.62353535
Registrant Phone Ext: 213
Registrant Fax:
Registrant Fax Ext:
Registrant Email: Luther.Torvalds@outlook.com
```

Figure 1.1: All the information is visible if WHOISGuard is ignored

However, it turns out that the guard isn't perfect, and all we had to do was find a WHOIS search engine that ignores the guard. After some digging, we found one, exposing all the juicy details of the registrant.

After getting the email address, we just had to send an email, and we were replied with the flag.

1.2 Flag

The flag for this challenge was `$CDDC19${IS_IT_I_AM_FAMOUS_NAO}`.

2 [R-1] Travel to the Past

The challenge heavily hinted that we would need a way to look at an older version of the website to find the flag, and that's exactly what we did.

2.1 Solution

There was a single snapshot of the website on the Wayback Machine¹ on the 20th of May, and browsing that snapshot revealed the flag on the homepage of a blog:

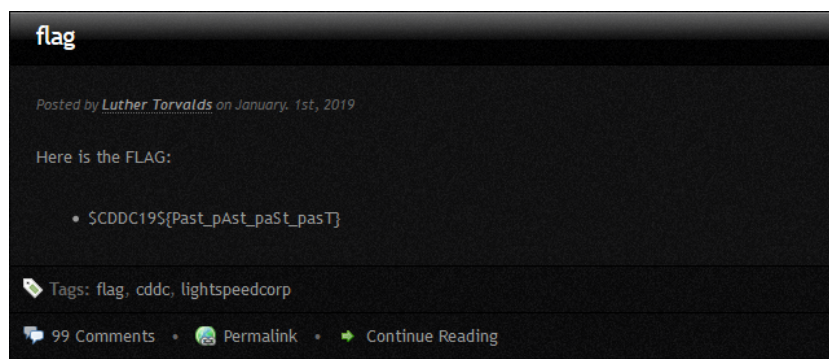


Figure 2.1: The flag is clearly visible

2.2 Flag

The flag for this challenge was `$CDDC19${Past_pAst_paSt_pasT}`.

[1] <https://archive.org/web/>

3 [R-2] I'm Sho Done With This

3.1 Solution

The hint was in the question title itself — turns out that there is a search engine called ‘ShoDan’², which claims to be the world’s first search engine for internet connected devices.

Sure enough, searching for our favourite company ‘lightspeedcorp’ yielded the following results:



Figure 3.1: The flag is reversed, for some reason

A trivial string-reverse later, the flag is obtained.

3.2 Flag

The flag for this challenge was \$CDDC19\${B1g-T3ch-Br0thers-4re-Wa7ching-U}.

[2] <https://www.shodan.io>

Programming

4 Count 1: Baby

This was a fairly simple code golfing challenge; inspection of the provided `count1-baby.py` file gave the character limit as 53 characters. The starting code is given as this:

```
1  #include <stdio.h>
2
3  int main(int argc, char *argv[])
4  {
5      int i;
6
7      for( i = 1 ; i < 10000 ; i++ )
8      {
9          printf("%d,", i);
10     }
11
12     return 0;
13 }
```

4.1 Solution

Compiling and running the program shows that the expected output of the minified program are the numbers 1 to 9999, separated by commas (with a trailing comma after 9999).

Besides the obvious steps like removing whitespace and newlines, knowledge of the C standard and ignorance of compiler warnings can yield the following transformations:

- `printf` is an implicitly defined standard library function; `stdio.h` can be removed
- Functions do not require a type specifier; `int main` can be shortened to `main`
- `main` does not need to take arguments, leaving `main()`
- The verifier rejects spaces, so `int i` has to be changed to `int(i)` (which is valid)
- `main` does not need to explicitly return 0

This yields the final program, at 53 characters long:

```
1 | main(){int(i);for(i=1;i<10000;i++){printf("%d",i);}}
```

Listing 4.1: The solution for Count 1: Baby

4.2 Flag

The flag for this challenge was `$CDDC19${Count2_is_waiting_Please_enjoy!}`.

5 Count 2: Wildness

The required output for this challenge is the same as the previous challenge (Count 1: Baby), but the restrictions have been tightened. This time, OP's friend 'doesn't like to go wild', or, in other words, the following characters cannot be used in the source code: `<== wlld ==>`. Additionally, the character limit has been reduced to 41.

5.1 Solution

Starting from the code for Count 1 (Listing 4.1), we can apply some more non-obvious tricks:

- Variables don't need a type specifier, and will default to `int`
- Variables with static storage duration will be initialised to 0 (obviating the need for `=`)
- `!=` is equivalent to a bitwise XOR (^) for integers (removing `<`)
- Taking advantage of `i++` semantics can replace `1e4` with `9999` (since `1` can't be used)

This yields the final program, again at exactly 41 characters:

```
1 | i;main(){for(;i++^9999;printf("%d,",i));}
```

Listing 5.1: The solution for Count 2: Wildness

5.2 Flag

The flag for this challenge was `$CDDC19${This_really_helps_m3_a_lot}`

6 Count 4: Madness - Filter

For this challenge, the character limit is slightly relaxed to 44 characters, but OP's friend has a faulty keyboard now, and all lowercase characters except those in 'mad printf' ('a', 'd', 'f', 'i', 'm', 'n', 'p', 'r', and 't') cannot be used.

6.1 Solution

Given that all the looping constructs (for, while, and even goto) cannot be used due to having illegal characters, the only logical solution is recursion.

- It is not illegal to call main recursively
- When called with no arguments, argc is 1; so main(i) will initialise i to 1
- 1E4 is equally as valid as 1e4, which is identical to 10000
- The first arm of the ternary operator can be omitted: x?:y is valid.

We came up with two independent solutions to this problem:

```
1 | main(i){printf("%d,",i++);if(i<1E4)main(i);}
```

Listing 6.1: The solution for Count 2: Wildness

```
1 | i;main(){printf("%d",++i);i>=9999?:main();}
```

Listing 6.2: An alternative solution for Count 2: Wildness

In both cases, the code was exactly 44 characters long.

6.2 Flag

The flag was `$CDDC19${Main_might_be_just_a_function_but_it_is_really_special!}`.

Reverse Engineering

7 LSCVM: Immaculate Invasion

Presented with a login prompt, the logical conclusion was that the flag would be accessible upon successfully logging into the server.

7.1 Solution

The tool of choice for this challenge was *Cutter*³, an open-source reverse-engineering framework which performs disassembly, function analysis, and function/value renaming, among other things.

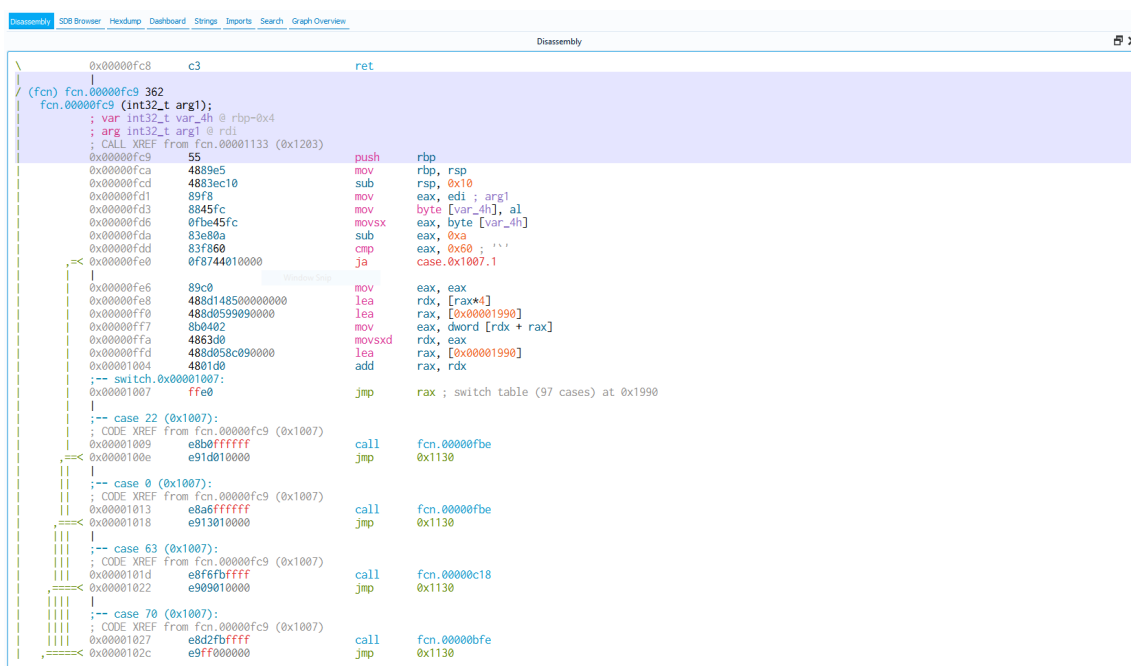


Figure 7.1: A screenshot of Cutter

The first hurdle was getting the program to run beyond its error message of [-] Flag file open error; upon inspection of the disassembly, the program tries to open a file named ‘flag’, which presumably contains the flag on the server-side — \$ touch flag convinced the program to start.

[3] <https://github.com/radareorg/cutter>

Given that the program asks for a login ID, the next step was to check for strings and `strcmps` in the code, and one was indeed found:

```
call    fcn.00001133
lea     rsi, str.lsc_user ; 0x2065 ; "lsc_user" ; const char *s2
lea     rdi, [0x00203140] ; 0x203140 ; const char *s1
call    sym.imp strcmp ; int strcmp(const char *s1, const char *s2)
test    eax, eax
je      0x14b3
```

Figure 7.2: A suspicious `strcmp`

Of course, no challenge is so easy, especially one that sat at above 950 points more than 24 hours after the qualifiers started. On further inspection, the operand of `strcmp` appeared to be the output of a call to `fcn.00001133`, which itself took some very long and suspicious strings as input: `eiMaAghMcAgjcMMdAgjcMMAAgjcMMAhhcMMdAijMaAcfMPPPPPPPP`.

Another discovery was this instruction: `cmp dword [var_1174h], 2`, where `var_1174h` contained the value of `argc`. Passing a second argument in the invocation (eg. `$./lscvm-ii x`) revealed a very useful debugging mode that printed what appeared to be a stack.

Following the call chain eventually led to the jump table for opcode dispatch, pictured in Figure 7.1. Some quick analysis revealed that this was a indeed stack-based virtual machine (*thankfully*). Combined with the debugging output, and the fact that the text output (the banner, login prompt, etc.) was printed using the VM instead of say `printf`, some basic opcodes were decoded:

Opcode	Pops	Pushes	Description
u to u	–	0 to 9	constant
A	a, b	a + b	add
M	a, b	a * b	multiply
P	a	–	<code>printf("%c", a)</code>

Piecing things together, it was inferred that we were most likely supposed to input a *program* when asked for the user ID, which would then print `lsc_user` as required. Given that ASCII values for the lowercase alphabets start at 97, a small program was quickly written that would take in a string and output opcodes that would print it (Listing 1.1).

After trying it out, however, we were quickly met with disappointment, as we were faced with our old friend `[-] Wrong id.`

After more digging, the user id was expected to be at the memory address 0x203140; using the renaming feature of Cutter, we could make it easier to spot:

```
call    fcn.00001133
lea     rsi, str.lsc_user ; 0x2065 ; "lsc_user" ; const char *s2
lea     rdi, var.mem_base.203140 ; 0x203140 ; const char *s1
call    sym.imp.strcmp ; int strcmp(const char *s1, const char *s2)
test    eax, eax
je      0x14b3
```

Figure 7.3: [0x203140] after being renamed into `var.mem_base.203140`

Scrolling through the opcode list while looking for the address (now renamed) yielded this very interesting function, mapped to opcode K:

```

(fcn) fcn.00000e22 86
    fcn.00000e22 ();
    ; var signed int var_4h @ rbp-0x4
    ; CALL XREF from fcn.00000fc9 (0x1107)
    0x00000e22      55                                push     rbp
    0x00000e23      4889e5                            mov      rbp, rsp
    0x00000e26      4883ec10                          sub      rsp, 0x10
    0x00000e2a      e89cfdffff                          call     fcn.00000bcb
    0x00000e2f      8945fc                            mov      dword [var_4h], eax
    0x00000e32      817dfc00ffffff                      cmp      dword [var_4h], 0xffffffff00
    0x00000e39      7c09                                jl       0xe44
    |
    0x00000e3b      817dfcff3f0000                      cmp      dword [var_4h], 0x3fff
    0x00000e42      7e16                                jle      0xe5a
    |
    ; CODE XREF from fcn.00000e22 (0xe39)
    -> 0x00000e44      8b45fc                            mov      eax, dword [var_4h]
    0x00000e47      89c6                                mov      esi, eax
    0x00000e49      488d3de80a0000                      lea      rdi, str.memory_write_access_violation_d ; 0x1938
    0x00000e50      b800000000                          mov      eax, 0
    0x00000e55      e876fbffff                          call     sym.imp.printf ; int printf(const char *format)
    |
    ; CODE XREF from fcn.00000e22 (0xe42)
    -> 0x00000e5a      e86cfdffff                          call     fcn.00000bcb
    0x00000e5f      89c1                                mov      ecx, eax
    0x00000e61      8b45fc                            mov      eax, dword [var_4h]
    0x00000e64      4863d0                            movsxd   rdx, eax
    0x00000e67      488d05d2222000                      lea      rax, var.mem_base.203140 ; 0x203140
    0x00000e6e      880c02                            mov      byte [rdx + rax], cl
    0x00000e71      b800000000                          mov      eax, 0
    0x00000e76      c9                                leave
    0x00000e77      c3                                ret

```

Figure 7.4: A function to surpass metal-gear write to memory?

There was also an accompanying opcode **E**, which read a value from memory. The memory buffer appeared to be an array of 32-bit values, and opcodes **K** and **E** took an index into this array. Slight modifications were made to `lscvm-deasciinator` to yield `lscvm-memoryinator` (basically keeping track of an offset and replacing ‘P’ with ‘K’), which would generate opcodes to write a string into a given offset in memory.

Finally, then, the challenge could be solved (the password was also in plain sight, hi_darkspeed-corp!):

```
$ ./lscvm-memoryinator
string: lsc_user
address: 0

ggdMMaKfgAfMcMfAbKjjMjAjAcKgfdMMfAdKfgAfMcMhAeKfgAfMcMfAfKcf
Mc fMMbAgKfgAfMcMeAhK

string: hi_darkspeed-corp!
address: 0

jeAiMaKhdfMMbKgfdMMfAcKcfMc fMMdKjjMjAhAeKfgAfMcMeAfKhdfMMcAg
KfgAfMcMfAhKfgAfMcMcAiKcfMc fMMbAjKcfMc fMMbAcfMKcfMc fMMbc fMMb
AKfddMMbc fMMcAKjjMjAjAbcfMMdAKfgAfMcMbAbcfMMeAKfgAfMcMeAbcfM
MfAKfgAfMcMcAbcfMMgAKfgAdMbc fMMhAK
^C

$ nc lscvm-ii.cddc19q.ctf.sg 9001

=== Welcome to LSCVM(LightSpeed Corp Virtual Machine) ===
ID : ggd...AhK
Password : jeA...hAK

Login Successful! $CDDC19${IcY_GrE37ings_Fr0M_LigHT5pEeDC0Rp}

lsc_user, Good Bye!
```

7.2 Flag

The flag for this challenge was `$CDDC19${IcY_GrE37ings_Fr0M_LigHT5pEeDC0Rp}`.

7.3 Further Analysis

While we were unable to actually come up with a quine to submit for the next LSCVM challenge (Quintessential Harlequin), we continued to take apart the VM (on the advice that it would be used again in the Finals!), since solving `lscvm-ii` did not require all of the opcodes (far from it).

Again, the renaming feature of Cutter was extremely helpful, and we were able to discover the address of the program counter, the base address of the stack, and (in the case of `lscvm-qh`), the mirror of stdout. Each buffer appears to be identically constructed with a fixed size of `0x4e200` bytes (`0x13880` 32-bit words), and the number of items stored after the last element (ie. at offset `0x4e200` from the base address).

We (eventually) managed to decipher all of the opcodes and their purpose:

Opcode	Pops	Pushes	Description
a to j	–	0 to 9	constant
A	a, b	a + b	add
B	–		stop execution immediately
C	x	–	call (jump to absolute instruction x)
D	x	–	pop (drop)
E	addr	value	read memory from addr
F	ofs	value	fetch from stack (ofs elms below top)
G	ofs	–	relative jump forward
H	ofs	value	same as F , but removes the element
I	x	–	printf("%d", x)
J	a, b	cmp	–1 if $a < b$, 0 if $a = b$, 1 if $a > b$
K	val, addr	–	writes val to memory at addr
M	a, b	a * b	multiply
P	x	–	printf("%c", x)
R	–	–	return
S	a, b	a - b	subtract
V	a, b	a / b	integer divide
Z	cond, ofs	–	jump (relative) if cond is 0

Of note are the **C** and **R** opcodes, which *call* and *return* respectively. There is another array which is only accessed by these instructions that functions as a *callstack*. **C** pushes the current program counter (*PC*) to this callstack, and **R** pops the return address from the callstack, and sets *PC* to it, moving execution back to the callsite.

Miscellaneous

8 Polyglot

8.1 Solution

The obvious solution was to put each sentence into Google Translate⁴; all 10 sentences were some variation of ‘the first letter of this language is the flag’. Coupled with the input format, given as [01][02][03][04][05][06]&[07][08][09][10]!, we were able to decode the flag.

Number	Language	Sentence
01	Hindi	Karakter pertama bahasa ini yang mengibarkan bendera.
02	Indonesian	
03	Chinese	Het eerste teken van deze taal vormt de vlag.
04	Dutch	
05	Danish	Det første tegn på dette sprog udgør flag.
06	Catalan	
07	Norwegian	El primer caràcter d'aquest idioma constitueix la bandera.
08	Spanish	
09	Hmong	Det første tegnet av dette språket utgjør flagget.
10	Croatian	

Assembled, the message was HIČ DDC&NSHC; it helped that there was a coherent message to verify that Google Translate didn't misdetect any of the languages.

8.2 Flag

The flag for this challenge was `$CDDC19${HIČ DDC&NSHC}`.

[4] <https://translate.google.com>

9 Super Strong TeleVision

With the obvious hints in the challenge statement, we got to work decoding.

SSTV, or Slow Scan Television, is a picture transmission method used mainly by amateur radio operators to transmit and receive static pictures via radio in monochrome or colour⁵. It can also be used to hide easter⁶ eggs⁷...

9.1 Solution

Using *PulseAudio* was used as a link between *VLC* and *QSSTV*⁸ (an open-source Linux SSTV application), we were able to decode the SSTV image.



Figure 9.1: QSSTV decoding the image; its frequency spectrum can be seen on the right

[5] en.wikipedia.org/wiki/Slow-scan_television

[6] https://wiki.kerbalspaceprogram.com/wiki/List_of_easter_eggs

[7] https://half-life.fandom.com/wiki/Portal_ARG

[8] <http://users.telenet.be/on4qz/qsstv/index.html>

In this case, *PulseAudio* acts as the pipe that redirects the .wav file as an input into the decoder. To achieve that, we first set up a virtual *null sink* and configure the system in the following manner:




Figure 9.2: Data Pipeline Setup

9.2 Flag

After some suspense as the SSTV audio played back in real time, the image was fully decoded, revealing the flag, which was `$CDDC19${Light$peepCorp-$$TV}`.



Figure 9.3: Decoded SSTV Image

Appendix

1 Code Listings

1.1 lscvm-deasciinator

```
1 // lscvm-deasciinator.cpp
2
3 #include <stdio.h>
4
5 #include <map>
6 #include <vector>
7 #include <string>
8 #include <iostream>
9
10 int main()
11 {
12     std::map<char, std::string> lookup;
13
14     lookup['a'] = "jjMjAhA";
15     lookup['b'] = "jjMjAiA";
16     lookup['c'] = "jjMjAjA";
17     lookup['d'] = "cfMcFMM";
18     lookup['e'] = "cfMcFMMbA";
19     lookup['f'] = "jiAcMM";
20     lookup['g'] = "jiAcMMbA";
21     lookup['h'] = "jeAiM";
22     lookup['i'] = "hdfMM";
23     lookup['j'] = "hdfMMbA";
24     lookup['k'] = "hdfMMcA";
25     lookup['l'] = "ggdMM";
26     lookup['m'] = "ggdMMbA";
27     lookup['n'] = "fgAfMcM";
28     lookup['o'] = "fgAfMcMbA";
29     lookup['p'] = "fgAfMcMcA";
30     lookup['q'] = "fgAfMcMdA";
31     lookup['r'] = "fgAfMcMeA";
32     lookup['s'] = "fgAfMcMfA";
```

```

33     lookup['t'] = "fgAfMcMgA";
34     lookup['u'] = "fgAfMcMhA";
35     lookup['v'] = "fgAfMcMiA";
36     lookup['w'] = "fgAfMcMjA";
37     lookup['x'] = "gcfcMMM";
38     lookup['y'] = "fgAfgAM";
39     lookup['z'] = "fgAfgAMB";
40     lookup['_'] = "gfdMMfA";
41     lookup['!'] = "fgAdM";
42     lookup['-'] = "fddMM";
43
44
45     std::string input;
46
47     while(true)
48     {
49         std::getline(std::cin, input);
50
51         std::string output;
52         for(size_t i = input.size(); i-- > 0; )
53             output += lookup[input[i]];
54
55         for(size_t i = 0; i < input.size(); i++)
56             output += "P";
57
58         printf("\n%s\n", output.c_str());
59     }
60 }

```

Listing 1.1: lscvm-deasciinator.cpp